


capacity

Call volume funnel




ets judged

Hash LSH

LEXICAL NOVELTY DESIGN

wording produces similar fingerprints. An unmatched fingerprint is treated as novel.



“reconcile invoice”
session C

NO CLOSE MATCH

C

PRIORITIZE

Similar words group together

Unfamiliar wording gets priority

lexical, not semantic

mate locally

labels -> E IDW estimate



480
CONCEPTS

5 MATCHED BUDGETS

Metadata tags

Fields used for identity, analysis, and retrospective scoring

IDENTITY

AGENT
contract-clause-checker

CHANNEL
m365-copilot-chat

RETROSPECTIVE EVALUATION METADATA

CORPUS
historical_300

DOMAIN
legal

TASK
compare indemnity + liability clauses

ended.

."

DIFFICULTY

borderline

SESSIONIZATION

inactivity

TURNS / TOOLS

2 / 2

PACKET

SPANS

5

DURATION

6 s

EMITTED TOKENS

490

TRUNCATED

false

RETROSPECTIVE ONLY • NOT SELECTION INPUT

expected label = 1 · pass

PASS

dth

Concept coverage


How much of the population’s behavioral variety did the selected sessions touch?

COUNT DISTINCT BEHAVIOR REGIONS TOUCHED



concept key = corpus | domain | task | difficulty




touched




touched





y won the first




GAP AT 50% BUDGET

+19.4 pts

Embedding coverage over Random at matched budget.

But: Random had lower selected-only pass-rate error in 4 of 5 budgets.




NOW Deterministic Random -> NEXT Semantic selection + IDW research

IDW is separate from selection and v4 IDW results were mixed.

me decision

state?



~~shared memory~~

~~Redis coordination~~


08 / 12


selection

izes?

pecialized workflows

rare behaviors






09 / 12

estimate

changes.

same L Changed: weighted donor angle


FAR DONORS




large weighted angle

geometry contributes

u_Lip = 20 pts



t without



5 traffic schedules


5 judge budgets

unlock - not a

keeper validation?

~119x

SMALLER STORED REPRESENTATION • 0.129 VS 15.36 MB



ma + compliance · production scale



THEN
ESP pilot: store 8 components and compute similarity directly